Minimax Iterative Dynamic Game

Olalekan Ogunmolu

Smilow Center for Translational Research,
The University of Pennsylvania, Philadelphia, PA

Sept. 11, 2019

The robustness conundrum

- How to know *a priori* a policy's robustness limits?

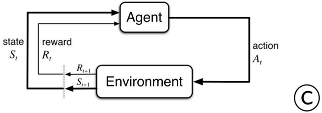

- How to inculcate robustness into multistage decision policies?

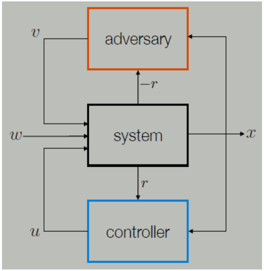

Problem Setup

- To quantify the brittleness, we optimize the stage cost

$$\max_{\mathbf{v}_t \sim \psi \in \Psi} \left[\sum_{t=0}^{T} \underbrace{c(\mathbf{x}_t, \mathbf{u}_t)}_{\text{nominal}} - \gamma \underbrace{g(\mathbf{v}_t)}_{\text{adversarial}} \right]$$

- To mitigate lack of robustness, we optimize the *cost-to-go*

$$\mathcal{J}_t(\mathbf{x}_t, \pi, \psi) = \min_{\mathbf{u}_t \sim \boldsymbol{\pi}} \max_{\mathbf{v}_t \sim \psi} \left(\sum_{t=0}^{T-1} \ell_t(\mathbf{x}_t, \mathbf{u}_t, \mathbf{v}_t) + L_T(\mathbf{x}_T) \right),$$

- and seek a saddle point equilibrium policy that satisfies

$$\mathcal{J}_t(\mathbf{x}_t, \pi^*, \psi) \leq \mathcal{J}_t(\mathbf{x}_t, \pi^*, \psi^*) \leq \mathcal{J}_t(\mathbf{x}_t, \pi, \psi^*),$$

Results: Brittleness Quantification

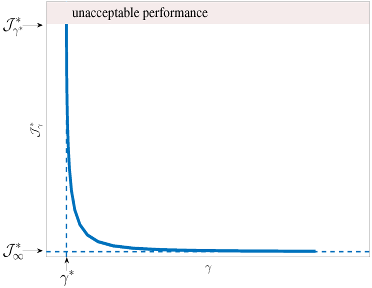

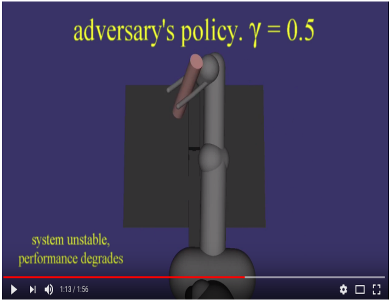

ILQG Algorithm Example

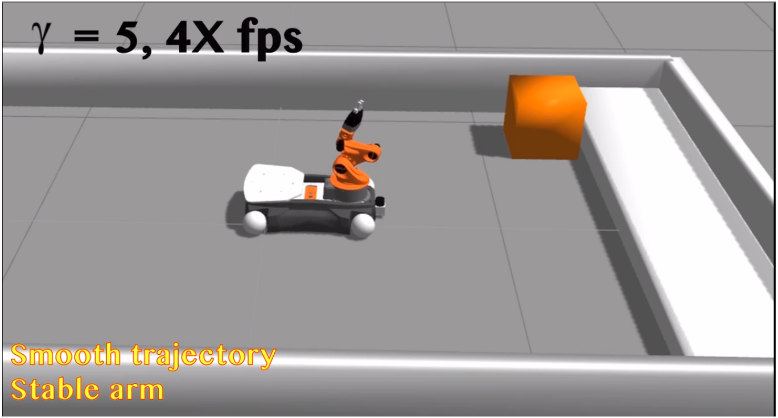

Results: Iterative Dynamic Game

$\mathbf{x}_1^\star$

$\mathbf{x}_2^\star$

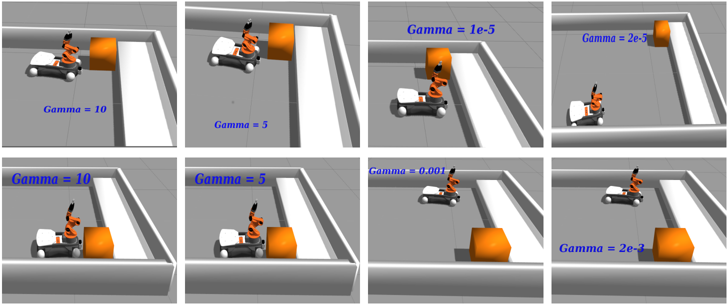

Table: *

End pose of the KUKA platform with our iDG formulation given different goal states and γ-values